

REST & GraphQL APIs

Table of Contents

- [Table of Contents](#)
- [Overview](#)
 - [Scope](#)
 - [Audience](#)
 - [Outcomes](#)
- [Core Principles](#)
- [API-First Process](#)
 - [Lifecycle](#)
 - [Artifacts](#)
- [REST API Design Requirements](#)
 - [3.1 Resource Modeling & URLs](#)
 - [3.2 Methods & Idempotency](#)
 - [3.3 Status Codes](#)
 - [3.4 Pagination, Filtering, Sorting](#)
 - [3.5 Caching](#)
 - [3.6 Error Envelope \(uniform\)](#)
 - [3.7 Security](#)
- [GraphQL API Design Requirements](#)
 - [4.1 Schema Modeling](#)
 - [4.2 Error Handling & Result Shapes](#)
 - [4.3 Performance](#)
 - [4.4 Authorization](#)
 - [4.5 Schema Evolution](#)
- [OpenAPI Integration \(REST\)](#)
- [Framework-Specific Guidance](#)
 - [6.1 Spring Boot](#)
 - [6.2 Next.js \(App Router\)](#)
- [REST + GraphQL Status & Error Mapping](#)
- [When to Use REST vs GraphQL](#)
- [Gateway & Edge](#)
- [Testing, Governance & CI/CD](#)
 - [Retry Policy](#)
 - [Purpose](#)
 - [Scope](#)
 - [Policy](#)
 - [Strategy](#)
 - [Example \(Spring Boot using Resilience4j\)](#)
 - [Example \(Next.js fetch wrapper\)](#)
 - [Observability](#)

Status	REVIEW
Version	1.0.0
Author(s)	@Daniel Steinhausf
Date of creation	Dec 18, 2025
Last updated	Dec 18, 2025
Approved on	
<name of the person(s) in charge to approve the document>	Date of approval: yyyy-mm-dd

Overview

This guideline establishes **consistent, scalable standards** for designing, implementing, and operating **REST and GraphQL APIs** across our engineering organization. It enforces an **API-First** approach, integrates **OpenAPI** for REST contracts, and provides framework-specific guidance for **Spring Boot (≥ 3.3)** and **Next.js**.

Scope

- **Design:** Resource modeling, schema design, versioning, error handling, pagination, caching, security, and retry policies.
- **Contracts:** OpenAPI (REST) and GraphQL schemas; mock servers; governance; CI validation.
- **Implementation:** Spring Boot services, Next.js consumers, and gateway concerns.
- **Operations:** Observability, reliability, performance, deprecation, and lifecycle management.

Audience

Backend and frontend engineers, architects, and platform teams working on services and clients that expose or consume REST/GraphQL APIs.

Outcomes

- Predictable APIs with **strong contracts, consistent error formats, robust security, and observable behavior.**
- Faster delivery via **contract-first development, mocking, and generated SDKs.**
- Reduced regressions through **linting, breaking-change checks, and contract tests.**

Core Principles

- **API-First:** Author and review API contracts **before** implementation (OpenAPI for REST, schema for GraphQL).
- **Consistency:** Shared conventions for naming, versioning, errors, pagination, authentication, and headers.
- **Backward Compatibility:** Prefer additive changes; deprecate before removing; publish sunset timelines.
- **Observability:** Tracing (OpenTelemetry), structured JSON logs with correlation IDs, per-endpoint/resolver metrics.
- **Security:** OIDC/JWT, scopes/roles, least privilege, strict validation, deny-by-default CORS, rate limiting.
- **Performance & Resilience:** Caching, idempotency keys, retries with backoff, circuit breakers, bulkheads, and efficient pagination.

API-First Process

Lifecycle

1. **Model** domain resources (REST) and entities/relations (GraphQL).
2. **Design** contracts: `openapi.yaml` and `schema.graphqls`.
3. **Mock** early; iterate with frontend consumers.
4. **Validate** in CI (lint, breaking-change checks, examples, security).
5. **Implement** behind contracts; enforce compatibility via contract tests.
6. **Publish** dev portal pages and SDKs; document changelogs.
7. **Operate** with SLOs, error budgets, and deprecation timelines.

Artifacts

- REST: `openapi.yaml` (+ error catalog, examples, deprecation log).
- GraphQL: `schema.graphqls` (+ directives for auth, deprecation notes, persisted queries list).

REST API Design Requirements

3.1 Resource Modeling & URLs

- Plural nouns for collections; singular for instances:

/v1/customers/{customerId}/orders/{orderId}

- Avoid verbs in paths; use HTTP methods for actions.

- Use natural subresources; otherwise link via IDs.

3.2 Methods & Idempotency

- Use `GET`, `POST`, `PUT`, `PATCH`, `DELETE` with standard semantics.

- **Idempotency-Key:** Required for POST actions with side effects (recommend 24h scope per operation+actor).

3.3 Status Codes

Use standard HTTP status codes consistently across all REST endpoints:

- `2xx` for success (`200 OK`, `201 Created`, `202 Accepted`, `204 No Content`).
- `3xx` for redirects (`301 Moved Permanently`, `302 Found`, `304 Not Modified`).
- `4xx` for client errors (`400 Bad Request`, `401 Unauthorized`, `403 Forbidden`, `404 Not Found`, `409 Conflict`, `422 Unprocessable Entity`, `429 Too Many Requests`).
- `5xx` for server errors (`500 Internal Server Error`, `502 Bad Gateway`, `503 Service Unavailable`, `504 Gateway Timeout`).

Each status code should align with its semantic meaning (e.g., `201` for resource creation, `202` for async processing, `409` for state conflicts).

Error responses must use the standardized envelope described in [REST & GraphQL APIs | 3.6 Error Envelope \(uniform\)](#).

For detailed mappings, GraphQL-specific patterns, and recommended headers, see: [REST & GraphQL APIs | REST + GraphQL Status & Error Mapping](#)

3.4 Pagination, Filtering, Sorting

- Offset/limit for simple lists; `cursor` for large/volatile datasets.
- Query params for filters (`?status=active`) and sorts (`?sort=-createdAt, name`).
- Multi-value filters allowed: `?tag=blue&tag=large`.

3.5 Caching

- `ETag` / `If-None-Match`, `Last-Modified` / `If-Modified-Since`.
- `Cache-Control` on GET when safe; leverage CDN for public cacheable endpoints.
- Consider `Prefer: return=minimal` for large payloads.

3.6 Error Envelope (uniform)

Use an RFC7807-inspired envelope across all REST services:

```

1  {
2    "error": {
3      "type": "https://errors.example.com/not-found",
4      "title": "Resource Not Found",
5      "status": 404,
6      "code": "CUSTOMER_NOT_FOUND",
7      "detail": "Customer 91e3... does not exist",
8      "instance": "traceId-12345",
9      "errors": [{ "field": "customerId", "message": "Invalid UUID" }]
10   }
11 }
```

3.7 Security

- OAuth2/OIDC with JWT; scopes per resource/action.
- mTLS for service-to-service when applicable.
- Strict CORS (deny-by-default); allow exact origins/headers/methods.

GraphQL API Design Requirements

4.1 Schema Modeling

- Keep mutations focused; use input types for complex args.
- Use Relay-style connections for pagination (`edges` , `pageInfo`).
- Prefer nullability transitions (add as nullable → migrate → enforce non-null when safe).

4.2 Error Handling & Result Shapes

- Most application errors return `200 OK` with `errors[]`; include `extensions.code`.
- Recommended mutation pattern with `userErrors`:

```

1  type Mutation {
2    createCustomer(input: CreateCustomerInput!): CreateCustomerPayload!
3  }
4
5  type CreateCustomerPayload {
6    customer: Customer
7    userErrors: [UserError!]!
8  }
9
10 type UserError {
11   field: String
12   message: String!
13   code: String
14 }
```

- Common `extensions.code` : `BAD_USER_INPUT` , `UNAUTHENTICATED` , `FORBIDDEN` , `NOT_FOUND` , `CONFLICT` , `PRECONDITION_FAILED` , `RATE_LIMITED` , `INTERNAL_SERVER_ERROR` .

Transport vs application errors: Transport-level failures (missing/invalid token, gateway timeouts) should use HTTP error codes; resolver/business errors remain `200` . See §7.

4.3 Performance

- **DataLoader** per request for batching & caching.
- Enforce max query depth/complexity; consider persisted queries & allow GET for CDN caching (`304` friendly).
- Avoid N+1 and deep nesting in resolvers.

4.4 Authorization

- Field-level auth when needed; enforce at resolvers (server-side), never only in the client.
- Map auth failures to `UNAUTHENTICATED` / `FORBIDDEN` in `extensions.code`.

4.5 Schema Evolution

- Deprecate fields/types with reasons:

```
1 type Customer {  
2   legacyId: ID @deprecated(reason: "Use id")  
3 }
```

- For breaking changes, route via the new schema/version behind the gateway if unavoidable.

OpenAPI Integration (REST)

- **Source of truth:** `openapi.yaml` at repo root with tags, reusable `components`, examples, and security schemes.
- **Docs:** Render with Swagger UI / Redoc; publish to dev portal.
- **Validation:** Lint with Spectral; detect breaking changes with `openapi-diff`.
- **SDKs:** Generate TypeScript and Java clients; publish to internal registries.

Minimal skeleton:

```
1 openapi: 3.0.3  
2 info: { title: Customers API, version: 1.0.0 }  
3 servers: [ { url: https://api.example.com/v1 } ]  
4 paths:  
5   /customers:  
6     get:  
7       tags: [Customers]  
8       summary: List customers  
9       parameters:  
10         - $ref: '#/components/parameters/limit'  
11         - $ref: '#/components/parameters/offset'  
12       responses:  
13         '200':  
14           description: OK  
15           content:  
16             application/json:  
17               schema: { $ref: '#/components/schemas/CustomerList' }  
18 components:  
19   parameters:  
20     limit: { name: limit, in: query, schema: { type: integer, minimum: 1, maximum: 100 } }  
21     offset: { name: offset, in: query, schema: { type: integer, minimum: 0 } }  
22   schemas:  
23     CustomerList:  
24       type: object  
25       properties:  
26         items: { type: array, items: { $ref: '#/components/schemas/Customer' } }  
27         totalCount: { type: integer, minimum: 0 }
```

Framework-Specific Guidance

6.1 Spring Boot

- **OpenAPI:** Use `springdoc-openapi` to generate and serve docs; configure OAuth2 security schemes.
- **GraphQL:** Use `spring-graphql`. Place schema at
`src/main/resources/graphql/schema.graphqls`.

- **Security:** Spring Security 6; JWT resource server (`.oauth2ResourceServer(jwt())`), method security (`@PreAuthorize`).
- **Observability:** Micrometer + OpenTelemetry exporter; include `traceId` / `spanId` in MDC/logs.
- **Validation:** Jakarta Bean Validation and global `@ControllerAdvice` / `DataFetcherExceptionResolver` to normalize errors.

REST example (controller and security)

```

1  @RestController
2  @RequestMapping("/v1/customers")
3  class CustomerController {
4      private final CustomerService service;
5
6      @GetMapping
7      public ResponseEntity<CustomerList> list(
8          @RequestParam(defaultValue = "50") int limit,
9          @RequestParam(defaultValue = "0") int offset) {
10         return ResponseEntity.ok(service.list(limit, offset));
11     }
12
13     @PostMapping
14     @PreAuthorize("hasAuthority('SCOPE_customers:write')")
15     public ResponseEntity<Customer> create(@Valid @RequestBody
CreateCustomerInput input) {
16         var created = service.create(input);
17         return ResponseEntity.status(HttpStatus.CREATED).body(created);
18     }
19 }
20
21 @Bean
22 SecurityFilterChain filterChain(HttpSecurity http) throws Exception {
23     http.authorizeHttpRequests(auth -> auth
24         .requestMatchers(HttpServletRequest.GET,
"/v1/customers/**").hasAuthority("SCOPE_customers:read")
25         .anyRequest().authenticated())
26         .oauth2ResourceServer(oauth -> oauth.jwt());
27     return http.build();
28 }
```

GraphQL error mapping (centralized)

```

1  @Component
2  class GraphQLErrorResolver implements DataFetcherExceptionResolver
{
3      @Override
4      public Mono<List<GraphQLError>> resolveException(Throwable ex,
DataFetchingEnvironment env) {
5          if (ex instanceof ConstraintViolationException cve) {
6              var error = GraphQLErrorBuilder.newError(env)
7                  .message("Invalid input")
8                  .errorType(ErrorType.ValidationError)
9                  .extensions(Map.of(
10                      "code", "BAD_USER_INPUT",
11                      "fieldErrors", cve.getConstraintViolations().stream()
12                          .map(v -> Map.of("field", v.getPropertyPath().toString(),
"message", v.getMessage()))
13                          .toList()
14                  ))
15                  .build();
16              return Mono.just(List.of(error));
17          }
18          var error = GraphQLErrorBuilder.newError(env)
19              .message("Unexpected error occurred")
20              .extensions(Map.of("code", "INTERNAL_SERVER_ERROR"))
21              .build();
22          return Mono.just(List.of(error));
23      }
24 }
```

DataLoader registry

```
1  @Bean
```

```

2 public DataLoaderRegistry dataLoaderRegistry(OrdersBatchLoader
ordersLoader) {
3     var registry = new DataLoaderRegistry();
4     registry.register("ordersByCustomer",
DataLoader.newDataLoader(ordersLoader));
5     return registry;
6 }

```

6.2 Next.js (App Router)

- **REST & GraphQL clients:**

- Generate **TypeScript** clients from OpenAPI and GraphQL Codegen from schema/operations.

- **Route handlers (app/api/.../route.ts):**

- Use for lightweight APIs/proxies; normalize transport errors (401/403/429/5xx).

- **Caching:**

- `cache: 'no-store'` for personalized or auth'd fetches; use `revalidate` /ISR for public cacheable GET.
- Persisted GraphQL queries via GET enable `ETag` / `304` at the edge/CDN.

- **Auth:**

- Middleware to attach/validate tokens; never expose secrets in client components.

Proxy example (edge runtime)

```

1 export const runtime = 'edge';
2
3 export async function GET(req: Request) {
4     const url = new URL(req.url);
5     const limit = url.searchParams.get('limit') ?? '50';
6
7     const res = await fetch(`.${process.env.API_BASE}/v1/customers?
limit=${limit}`, {
8         headers: { Authorization: `Bearer ${await getAccessToken()}` },
9         cache: 'no-store'
10    });
11
12    return new Response(await res.text(), { status: res.status, headers:
res.headers });
13 }

```

REST + GraphQL Status & Error Mapping

Guiding rule: In GraphQL, **transport-layer** failures (auth, gateway, rate limit, timeouts) should return appropriate **HTTP status codes**. **Application-layer** (resolver/business) errors typically return `200 OK` with entries in `errors[]` and a meaningful `extensions.code`.

HTTP Code	Name	REST Use Case	GraphQL (Transport Layer)	GraphQL (Application Layer) – recommended <code>extensions. code</code>	Recommended Headers / Notes

200	OK	Successful GET/PUT/PATCH/POST returning data.	Successful GraphQL response (may contain <code>data</code> and/or <code>errors</code>).	<code>BAD_USER_IN</code> <code>PUT</code> , <code>INTERNAL_SE</code> <code>RVER_ERROR</code> , domain codes (e.g., <code>CUSTOMER_NO</code> <code>T_FOUND</code>).	Use caching headers for safe queries.
201	Created	Resource created; include <code>Location</code> .	Rare (mutations usually return 200 with payload).	Use payload + <code>userErrors</code> for validation issues.	<code>Location</code> with new resource URI (REST).
202	Accepted	Async job started; poll status.	Often still 200 with the job handle in the payload.	Payload includes job status (<code>PENDING</code> , <code>IN_PROGRES</code> <code>S</code>).	Provide job ID or status endpoint.
204	No Content	Success without body (DELETE).	Avoid—GraphQL expects a body.	Use payload flag (e.g., <code>deleted: true</code>).	Prefer explicit confirmations.
301	Moved Permanently	Resource moved.	Rarely used for GraphQL endpoints.	N/A	If used, set <code>Location</code> .
302	Found	Temporary redirect.	Generally avoid for GraphQL.	N/A	Gateway only.
304	Not Modified	Conditional GET with <code>ETag</code> / <code>Last-Modified</code> .	GET persisted queries can use <code>ETag</code> and <code>If-None-Match</code> .	N/A	Enable CDN/edge caching.
400	Bad Request	Malformed request/validation.	Syntax/validation errors of the GraphQL request can result in a <code>400</code> .	Domain validation → <code>200</code> + <code>BAD_USER_IN</code> <code>PUT</code> .	Use consistent error envelopes (REST).

401	Unauthorized	Missing/invalid auth.	401 from gateway/resource server.	Resolver-level auth → 200 + UNAUTHENTICATED .	WWW-Authenticate header.
403	Forbidden	Authenticated but no permission.	403 at gateway/policy layer.	FORBIDDEN at resolver.	Include required scopes in detail.
404	Not Found	Resource/endpoint missing.	Unknown GraphQL endpoint → 404 .	Missing entity → 200 , data.field = null + NOT_FOUND (optional).	Keep the REST envelope consistent.
405	Method Not Allowed	HTTP method not supported.	Rare for GraphQL (POST/GET).	N/A	Include Allow header.
409	Conflict	State conflict or version mismatch.	Use transport 409 if the operation cannot start.	CONFLICT or domain-specific (e.g., VERSION_CONFLICT).	Provide remediation hints.
410	Gone	Resource permanently removed.	Rare; prefer app-level handling.	Explain deprecation in the payload.	Pair with Sunset headers in REST.
412	Precondition Failed	If-Match / If-Unmodified-Since failed.	Typically handled in REST.	PRECONDITION_FAILED in extensions.	Favor ETag concurrency in REST.
415	Unsupported Media Type	Wrong Content-Type .	415 if non-JSON or invalid GraphQL type.	N/A	Enforce application/json .
422	Unprocessable Entity	Semantic validation failure.	Schema-valid but semantically invalid → often 200 + BAD_USER_IN PUT .	BAD_USER_IN PUT with field errors.	Include per-field details.

429	Too Many Requests	Rate limit exceeded.	429 at gateway; GraphQL not executed.	Resolver throttling → 200 + RATE_LIMITED .	Retry-After (seconds or HTTP date).
500	Internal Server Error	Generic server error.	Server/runtime failure → 500 .	Resolver exception → 200 + INTERNAL_SERVER_ERROR .	Don't leak internals; log traceId.
502	Bad Gateway	Upstream invalid response.	Gateway/proxy to GraphQL returns 502 .	N/A	Trace upstream.
503	Service Unavailable	Maintenance/overload.	Gateway returns 503 .	Resolver-level overload → optionally 200 with error; prefer 503 if not processed.	Retry-After ; circuit breakers.
504	Gateway Timeout	Upstream timeout.	Gateway returns 504 .	N/A	Tune timeouts; propagate correlation IDs.

GraphQL error payload examples:

- Validation

```

1  {
2    "data": { "createCustomer": null },
3    "errors": [
4      {
5        "message": "Email is invalid",
6        "path": ["createCustomer"],
7        "extensions": {
8          "code": "BAD_USER_INPUT",
9          "fieldErrors": [{ "field": "email", "message": "Must be a valid
10            email address" }],
11          "traceId": "abc123"
12        }
13      }
14    ]
15  }

```

- Auth

```

1  {
2    "data": null,
3    "errors": [
4      {
5        "message": "Authentication token missing",
6        "extensions": { "code": "UNAUTHENTICATED", "traceId": "abc123" }
7      }
8    ]
8  }

```

- Not Found (nullable field)

```

1  {
2    "data": { "customer": null },

```

```

3   "errors": [
4     "message": "Customer not found",
5     "path": ["customer"],
6     "extensions": { "code": "NOT_FOUND", "traceId": "abc123" }
7   }
8 }
```

- **Internal Resolver Error**

```

1 {
2   "data": null,
3   "errors": [
4     "message": "Unexpected error occurred",
5     "extensions": { "code": "INTERNAL_SERVER_ERROR", "traceId": "abc123" }
6   ]
7 }
```

When to Use REST vs GraphQL

- **REST:** Stable resource boundaries; cacheable GETs; external/public APIs; strong OpenAPI ecosystem.
- **GraphQL:** Complex UIs; avoid over-/under-fetching; rich relationships; field-level auth; internal/front-end heavy apps.

Hybrid: Provide REST externally and GraphQL internally; unify via a gateway with common auth, observability, and rate limits.

Gateway & Edge

- Central API gateway for **auth, rate limiting, request/response size caps, header normalization, compression, error translation**, and routing (REST & GraphQL).
- **Caching:** CDN for REST GET; **persisted GraphQL queries** (GET) for edge caching and **304** handling.

Testing, Governance & CI/CD

- **Contract tests:** REST (Pact), GraphQL (breaking-change checks, schema linting, query whitelists).
- **Integration:** Testcontainers in Spring.
- **Performance:** k6/Gatling; track P95/P99 and error budgets.
- **Security:** Scope enforcement tests; ZAP fuzzing; JWT/CSRF/CORS checks.
- **Governance:** Spectral (OpenAPI), eslint-plugin-graphql; block merges on breaking changes or missing examples/security.
- **Automation:** Generate & publish SDKs on release; dev portal updates; deprecation trackers.

Retry Policy

Purpose

Ensure resilience and graceful degradation when backend services or upstream dependencies are temporarily unavailable.

Scope

Applies to:

- **Client-side SDKs** (generated from OpenAPI or GraphQL schema).
- **Service-to-service calls** within Spring Boot.
- **Next.js API routes** acting as proxies.

Policy

- **Trigger conditions:**

- Any **5xx HTTP status code** (500 , 502 , 503 , 504).
- **Connection timeouts** or transient network failures.

- **Excluded:**

- **4xx client errors** (do not retry).
- **Validation/business logic errors** (GraphQL BAD_USER_INPUT , REST 422).

Strategy

- **Max retries:** 3 attempts (initial + 2 retries).

- **Backoff:** Exponential with jitter:

- Initial delay: **500 ms**.
- Multiply by a factor of **2** for each retry.
- Add random jitter ±20% to avoid thundering herd.

- **Cap:** Maximum delay per attempt = **5 seconds**.

- **Timeout:** Overall request timeout = **15 seconds** (including retries).

- **Circuit breaker:**

- Open after **5 consecutive failures**.
- Half-open after **30 seconds**.

- **Idempotency:**

- Only retry **idempotent operations** (GET, HEAD, PUT, DELETE).
- For POST, require **Idempotency-Key** header.

Example (Spring Boot using Resilience4j)

```
1 @Bean
2 public RetryConfig retryConfig() {
3     return RetryConfig.custom()
4         .maxAttempts(3)
5         .waitDuration(Duration.ofMillis(500))
6         .intervalFunction(IntervalFunction.ofExponentialBackoff(500, 2.0))
7         .retryExceptions(IOException.class, TimeoutException.class)
8         .ignoreExceptions(HttpClientErrorException.class)
9         .build();
10 }
11
12 @Bean
13 public RetryRegistry retryRegistry() {
14     return RetryRegistry.of(retryConfig());
15 }
```

Example (Next.js fetch wrapper)

```
1 async function fetchWithRetry(url: string, options: RequestInit,
2   retries = 2, delay = 500): Promise<Response> {
3   try {
4     const res = await fetch(url, options);
5     if (res.ok || (res.status >= 400 && res.status < 500)) return res;
6     throw new Error(`Retryable status: ${res.status}`);
7   } catch (err) {
8     if (retries <= 0) throw err;
9     await new Promise(r => setTimeout(r, delay + Math.random() * delay
10 * 0.2));
11     return fetchWithRetry(url, options, retries - 1, delay * 2);
12   }
13 }
```

Observability

- Log each retry attempt with `traceId`, attempt count, and delay.
- Emit metrics: `retry_attempts_total`, `retry_failures_total`.