

Service Cron Execution

Table of Contents

- [Table of Contents](#)
- [Overview](#)
- [Concept
 - Execution behavior
 - Benefits of this architecture](#)
- [Database Requirements for Distributed Cron Execution
 - PostgreSQL functions used:
 - Lock Key Rules](#)
- [Environment Variable Schema for Cron Jobs
 - JSON schema
 - Example
 - Why this schema?](#)
- [Further Structuring for Large Organizations
 - Lock-key allocation strategy
 - Naming convention](#)
- [Spring Boot — Distributed Cron Execution \(AOP-Based\)
 - Annotation
 - Distributed Lock Aspect
 - Scheduler Registrar
 - Example Distributed Cron Job](#)
- [Node.js / Next.js Cron Execution
 - Distributed Cron Helper
 - Using CRON_JOBS](#)
- [Lock-Key Registry](#)
- [Summary](#)

Status	DRAFT
Version	1.0.0
Author(s)	@Daniel Steinhau f
Date of creation	Feb 17, 2026
Last updated	Feb 17, 2026
Approved on	<name of the person(s) in charge to approve the document>
	Date of approval: yyyy-mm-dd

Overview

This guideline defines the unified architecture for **distributed-safe cron execution** across backend services implemented in:

- **Spring Boot** (Java, Hibernate, Liquibase)
- **Node.js**, including **self-hosted Next.js backend runtimes**

Cron jobs must run **exactly once cluster-wide**, even when multiple instances of the same service are deployed (e.g., in Kubernetes).

To guarantee this, services use **PostgreSQL advisory locks**, ensuring that only **one** instance executes a scheduled job at any point in time.

Concept

Each application instance (Spring Boot / Node.js / Next.js server) locally triggers cron jobs.

Before executing a job, the instance attempts to acquire a distributed lock:

```
1 SELECT pg_try_advisory_lock(<lock_key>);
```

Execution behavior

Condition	Behavior
Lock acquired	Execute job
Lock NOT acquired	Skip execution

After execution:

```
1 SELECT pg_advisory_unlock(<lock_key>);
```

Benefits of this architecture

- No external infrastructure required (e.g., Redis, Zookeeper)
- Very fast (PostgreSQL in-memory locking)
- Fully distributed and language-agnostic
- Safe during rolling deployments and autoscaling
- No schema required

Database Requirements for Distributed Cron Execution

PostgreSQL advisory locks:

- **Do NOT require tables**
- **Do NOT write to disk**
- Are **session-scoped**
- Are automatically released if:
 - the process crashes
 - the DB connection drops

PostgreSQL functions used:

Function	Purpose
<code>pg_try_advisory_lock(key)</code>	Attempts to acquire lock without waiting
<code>pg_advisory_unlock(key)</code>	Releases the lock
Session close	Auto-releases all locks

Lock Key Rules

- Each cron job must have a **globally unique lock key**
- Keys must be consistent across environments (dev/stage/prod)
- Services must **not** reuse lock keys

Environment Variable Schema for Cron Jobs

To define cron jobs dynamically from environment configuration, each service declares:

```
1 CRON_JOBS=<JSON object defining per-job schedule + lock key>
```

JSON schema

Each job is defined as:

```
1 "<jobName>": {
2   "schedule": "<cron expression>",
3   "lockKey": <integer number>
4 }
```

Example

```
1 CRON_JOBS={
2   "tokenCleanup": { "schedule": "0 */1 * * * *", "lockKey": 5101 },
3   "sessionPrune": { "schedule": "0 */5 * * * *", "lockKey": 5102 }
4 }
```

Why this schema?

- One env var per service → easy to audit
- Job configuration is declarative
- Job schedule + key stored together
- Fully compatible with Node.js, Next.js, and Spring Boot
- Easy to validate at startup

Further Structuring for Large Organizations

Large systems with many microservices require **lock-key ranges** to avoid collisions.

Lock-key allocation strategy

Each service is assigned a **start key**.

Jobs increment keys sequentially:

```
1 <service_start> + <job_index>
```

This guarantees global uniqueness.

Example:

```
1 lt-credit-service → 1301-1399
2 lt-invoice-service → 1551-1599
```

Naming convention

To avoid ambiguity:

```
1 <service>.<jobName>
```

Example:

```
1 "lt-credit-service.dailyLimitRefresh": { "schedule": "...", "lockKey": 1301 }
```

Spring Boot — Distributed Cron Execution (AOP-Based)

Spring Boot integrates distributed cron execution *transparently* via:

- A custom annotation `@DistributedScheduled`
- A lock aspect that uses PostgreSQL advisory locks
- A scheduler registrar that discovers annotated methods

This avoids developers having to manually implement locking.

Annotation

```
1 import java.lang.annotation.*;
2
3 @Retention(RetentionPolicy.RUNTIME)
4 @Target(ElementType.METHOD)
5 public @interface DistributedScheduled {
6     String cron();
7     long lockKey();
8 }
```

Distributed Lock Aspect

```
1 @Slf4j
2 @Component
3 public class DistributedLockAspect {
4
5     private final JdbcTemplate jdbc;
6
7     public DistributedLockAspect(JdbcTemplate jdbc) {
8         this.jdbc = jdbc;
9     }
10
11    public void execute(Object bean, Method method, long lockKey) {
12        boolean acquired = false;
13
14        try {
15            acquired = Boolean.TRUE.equals(
16                jdbc.queryForObject("SELECT pg_try_advisory_lock(?)",
17                    Boolean.class, lockKey)
18            );
19
20            if (!acquired) {
21                log.info("Skipping {} - lock {} held elsewhere",
22                    method.getName(), lockKey);
23                return;
24            }
25
26            log.info("Executing {} with lock {}", method.getName(),
27                lockKey);
28            method.invoke(bean);
29
30        } catch (Exception e) {
31            log.error("Error executing job {}", method.getName(), e);
32
33        } finally {
34            if (acquired) {
35                try {
36                    jdbc.queryForObject("SELECT
37 pg_advisory_unlock(?)", Boolean.class, lockKey);
38                } catch (Exception unlockError) {
39                    log.error("Failed to unlock key {}", lockKey,
40                    unlockError);
41                }
42            }
43        }
44    }
45 }
```

Scheduler Registrar

```
1 @Component
2 public class DistributedSchedulerRegistrar implements
3 SchedulingConfigurer {
```

```

3
4     private final ApplicationContext ctx;
5     private final TaskScheduler scheduler;
6     private final DistributedLockAspect aspect;
7
8     public DistributedSchedulerRegistrar(ApplicationContext ctx,
9                                         TaskScheduler scheduler,
10                                        DistributedLockAspect aspect)
11    {
12        this.ctx = ctx;
13        this.scheduler = scheduler;
14        this.aspect = aspect;
15    }
16
17    @Override
18    public void configureTasks(ScheduledTaskRegistrar registrar) {
19        registrar.setTaskScheduler(scheduler);
20
21        ctx.getBeansWithAnnotation(Component.class).forEach((name,
22 bean) -> {
23            Class<?> targetClass = AopUtils.getTargetClass(bean);
24
25            for (Method m : targetClass.getMethods()) {
26                DistributedScheduled ann =
27                    m.getAnnotation(DistributedScheduled.class);
28
29                if (ann != null) {
30                    registrar.addCronTask(
31                        () -> aspect.execute(bean, m, ann.lockKey()),
32                        ann.cron()
33                    );
34                }
35            });
36        });
37    }
38 }

```

Example Distributed Cron Job

```

1  @Component
2  public class CleanupJob {
3
4      @DistributedScheduled(cron = "0 */1 * * * *", lockKey = 5001)
5      public void cleanup() {
6          System.out.println("Distributed cleanup executed");
7      }
8 }

```

Node.js / Next.js Cron Execution

Next.js (self-hosted) supports cron execution only in the **Node.js server runtime**:

- NOT in API routes
- NOT in Edge runtime
- NOT in browser React

Use the same distributed locking helper for all Node-based runtimes.

Distributed Cron Helper

```

1  const cron = require('node-cron');
2
3  class DistributedCron {
4      constructor(pgClient) {
5          this.pg = pgClient;
6      }
7
8      schedule(schedule, lockKey, job) {
9          cron.schedule(schedule, async () => {
10              let acquired = false;
11
12              try {
13                  const res = await this.pg.query(
14                      "SELECT pg_try_advisory_lock($1) AS acquired",

```

```

15         [lockKey]
16     );
17     acquired = !!res.rows?.[0]?.acquired;
18
19     if (!acquired) {
20       console.log(`[DistributedCron] skipping job ${lockKey}`);
21       return;
22     }
23
24     console.log(`[DistributedCron] executing job ${lockKey}`);
25     await job();
26
27   } finally {
28     if (acquired) {
29       await this.pg.query("SELECT pg_advisory_unlock($1)",
30     [lockKey]);
31   }
32 });
33 }
34 }
35
36 module.exports = DistributedCron;

```

Using CRON_JOBS

```

1 const CRON_JOBS = JSON.parse(process.env.CRON_JOBS);
2
3 const dcron = new DistributedCron(pg);
4
5 Object.entries(CRON_JOBS).forEach(([jobName, cfg]) => {
6   dcron.schedule(cfg.schedule, cfg.lockKey, async () => {
7     console.log(`Executing job: ${jobName}`);
8     // job implementation
9   });
10 });

```

Lock-Key Registry

Each service is assigned a starting key. The range/number of cron jobs is 50 as of now:

Service Name	Start Key
lt-carbone-connector	1
lt-compass-connector	51
lt-crif-connector	101
lt-hyundai-connector	151
lt-idnow-connector	201
lt-intrum-connector	251
lt-ksv-connector	301
lt-mesoneer-connector	351
lt-odoo-connector	401
lt-pdf4Me-connector	451
lt-postmark-connector	501
lt-schufa-connector	551
lt-tesla-connector	601

lt-zek-connector	651
lt-zefix-connector	701
bawag-proxy	751
calculator	801
dealer	851
dealer-service-v1	901
ebicsbatch	951
insurance-service-v1	1001
keycloak-theme	1051
leasing-calculator-service-v1	1101
lt-accounting-service	1151
lt-aml-service	1201
lt-company-registry-check-service	1251
lt-credit-service	1301
lt-customer-sync-service	1351
lt-flowapp-v2-service	1401
lt-invoice-delivery-service	1451
lt-invoice-payment-sync-service	1501
lt-invoice-service	1551
lt-invoice-updater-service	1601
lt-kyc-service	1651
lt-partner-service	1701
lt-payment-plan-service	1751
lt-proxy-service	1801
partner	1851
partner-service-v1	1901
proxy	1951
proxy-service-v1	2001
lt-crm-be	2051
lt-crm-fe	2101
lt-flowapp	2151

Summary

This guideline provides:

- A robust model for **distributed cron execution**
- Architecture compatible with Spring Boot, Node.js, and Next.js
- A standardized **CRON_JOBS JSON schema** defining job schedule + lock key
- A global lock-key strategy and registry
- Ready-to-use implementation for both runtimes